

Grok 4.6 in Cursor vs. a bare-bones harness

The first model-times-harness study on VulcanBench: one model, two delivery systems, on twenty-three real merged post-cutoff PRs. The bare-bones system is Report No. 14; the Cursor system is new, three attempts per task at four effort levels.

ABSTRACT

Cursor is worth 14.5 points at high effort. At every other level the two systems land within noise of each other. At the high setting, Grok 4.6 scores 88.4% inside Cursor and 73.9% through VulcanBench’s minimal reference loop, and that is the only level where the gap is larger than the error bars. Low is +1.5, xhigh +7.2, and at medium the reference loop is ahead by 2.9. The clearest difference is what happens above medium: the raw API drops from 87.0% to 73.9% as effort rises to high, while the same step inside Cursor goes up, from 84.1% to 88.4%. Inside Cursor the effort knob barely matters at all (84.1, 84.1, 88.4, 85.5). The agent also tried to reach the web 299 times across the 276 runs, more often at higher effort, and was refused every time.

TABLE 1. The two systems

	System A: bare-bones	System B: inside Cursor
Path	xAI API, VulcanBench minimal reference loop	cursor-agent CLI, Cursor’s own scaffold
Effort control	reasoning_effort sent per request	baked into the model id (cursor-grok-4.6-low ... -xhigh)
Observability	tokens, cache reads, dollars at list price	no token or cost reporting; web denied, workspace outside repo
Attempts	1 per task (Report No. 14, \$52.70 recorded)	3 per task, 23/23 tasks, 276 runs, 0 contaminated

TABLE 2. pass@1 by nominal effort level (mean per-task success rate; ±1 stderr)

Effort	In Cursor	Bare-bones	Delta	Cursor time/task	Bare-bones time/task
low	84.1% ±7.5	82.6% ±8.1	+1.5	3.5 min	4.8 min
medium	84.1% ±7.2	87.0% ±7.2	−2.9	4.6 min	14.2 min
high	88.4% ±6.5	73.9% ±9.4	+14.5	6.9 min	16.3 min
xhigh	85.5% ±7.2	78.3% ±8.8	+7.2	6.5 min	15.8 min

Cursor columns: 69 runs each (three attempts per task, 23/23 tasks, zero contaminated on either audit channel). Bare-bones columns: 23 runs each, single attempt (Report No. 14). Times are medians and include provider latency. Only the high-effort gap exceeds the combined uncertainty of its pair.

FINDINGS

- The advantage is one level wide.** Only the high setting produces a gap bigger than its error bars: +14.5 against ±6.5. Low is +1.5, xhigh is +7.2, and at medium the reference loop is ahead by 2.9. Across the whole knob the two systems average about four points apart.
- Cursor prevents the high-effort collapse.** Report No. 14 found the raw API falling from 87.0% at medium to 73.9% at high. Inside Cursor the same step rises, 84.1% to 88.4%. The scaffold helps at the setting where the model tends to overwork a problem, and only there.
- The effort knob barely matters inside Cursor.** 84.1, 84.1, 88.4, 85.5 is a 4.3-point spread, against 13.1 through the raw API. Which harness you use matters more than which effort you pick.
- Higher effort means more web attempts.** Refused fetch attempts climb from 29 at low to 66, 96, and then 108 at xhigh, 299 in total, every one denied. A benchmark that allowed browsing would flatter the high-effort settings most. Every figure here comes from runs where no fetch went through.
- The stubborn tasks differ by system.** `sqlglot-canonicalize-internal-names` is never solved by either system at any setting. `itertools-strip-prefix` goes the other way: every raw-API level solves it, and every Cursor attempt fails.

TABLE 3. Task-level consistency inside Cursor (contained runs, three attempts each)

Effort	Solved every attempt	Mixed	Never solved
low	19/23	2	2
medium	18/23	3	2
high	20/23	1	2
xhigh	19/23	2	2

The same two tasks go unsolved at every level: `itertools-strip-prefix` and `sqlglot-canonicalize-internal-names`.

CAVEATS

The high-effort gap could come from Cursor’s scaffold (context management, tool design, stopping rules), from how Cursor’s effort-suffixed model ids map to the API’s `reasoning_effort` values, or from different serving. “Cursor Grok 4.6” is Cursor’s own branding and may be a tuned or differently-hosted deployment; these are indistinguishable from outside, and per-request token counts, which Cursor does not report, would settle it. Attempt counts differ between systems (3 vs 1), so per-task flips carry more weight on the Cursor side. Bare-bones costs in Report No. 14 are lower bounds: xAI reports reasoning tokens outside `completion_tokens` and the harness under-counted them during that sweep (since fixed). The xhigh column was collected in two sessions days apart after promotional credits ran out mid-level; every other level ran in one continuous pass. The audit certifies that no run reached the web or the benchmark’s own files; it cannot certify what was in the model’s training data.

METHOD

VulcanBench v3: 23 tasks from real merged post-cutoff PRs (Python 9, TypeScript 4, Rust 4, Go 3, JavaScript 3). Both systems are graded by the same deterministic hidden tests in Docker; model judges disabled. System A ran 2026-08-12 (Report No. 14); System B ran 2026-08-15/16 with `cursor-agent 2026.08`, non-fast model ids, Cursor’s sandbox enabled, web tools denied through a workspace permissions file, and an agent workspace created outside the repository so that no task definition, gold patch, or hidden test exists above the agent’s working directory; VulcanBench setup and verification run in Docker over that workspace. Every run carries an integrity audit of both leak channels and all 276 are clean. `pass@1` is the mean per-task success rate, so uneven attempt counts do not bias it. This is the first VulcanBench study measuring one model through two harnesses; the series continues with other model-times-harness combinations.

